


a

F1a Pearson, three readings



log10 λ, none, mask, rand

b


F1b held-out loss



log10 λ

c

F1c edge recall



log10 λ

d


F1d divergence



log10 λ

e


F1e gradient ratio



probe epoch

f


F1f random control



target

g

F2a Pearson, three readings



mechanism (12 columns)

h


F2b held-out loss



mechanism

i


F2c edge recall



mechanism

j


F2d val Pearson by epoch



epoch

k


F2e train Pearson by epoch



epoch

l


F2f val loss by epoch



epoch

m


F3a 60 epochs, val Pearson



epoch (0-59)

n

F3b 60 epochs, val loss



epoch (0-59)

o


F3c discovery



top-k off-graph attention

p


F3d overruling



α_{ij} / (1/d_{ij}), binned

q

F3e divergence, all arms



arm

r

F3f seed table

