d **Reading the three fits: the question each one answers and what each outcome says to do next**

Is the loss still falling with D at the full pool, on the matched read?

yes	a new screen beats a larger model: a doubling of distinct records buys 2^-beta of the reducible loss, and beta with its interval is the number to carry
no	data saturates at this size: the curve has met the plateau E + A N^-alpha, and the next unit of effort goes to N or to the floor

Is the loss still falling with N at the full pool?

yes	capacity binds: the next size pays, and alpha says how much per doubling; check that lr was right for the top of the ladder first
no	the sizes sit on the plateau E + B D^-beta: more data before more parameters, and the N exponent is descriptive, not a forecast

Where does the fitted E sit against the replicate floor of each screen?

below	the fit is wrong: no predictor beats the replicate noise, so E and the exponent traded off on too few points
near	the model family is near the noise limit on that screen; the remaining loss is measurement, not modeling
far above	the model family is the limit, not the data: a different inductive bias before a bigger grid

gate on every fit: the largest point is held out and predicted from the rest. Predicted within the bootstrap interval, the exponent extrapolates one step; missed, it describes the points it was fit on and nothing beyond them. The floored form needs four distinct x in sample for this test.