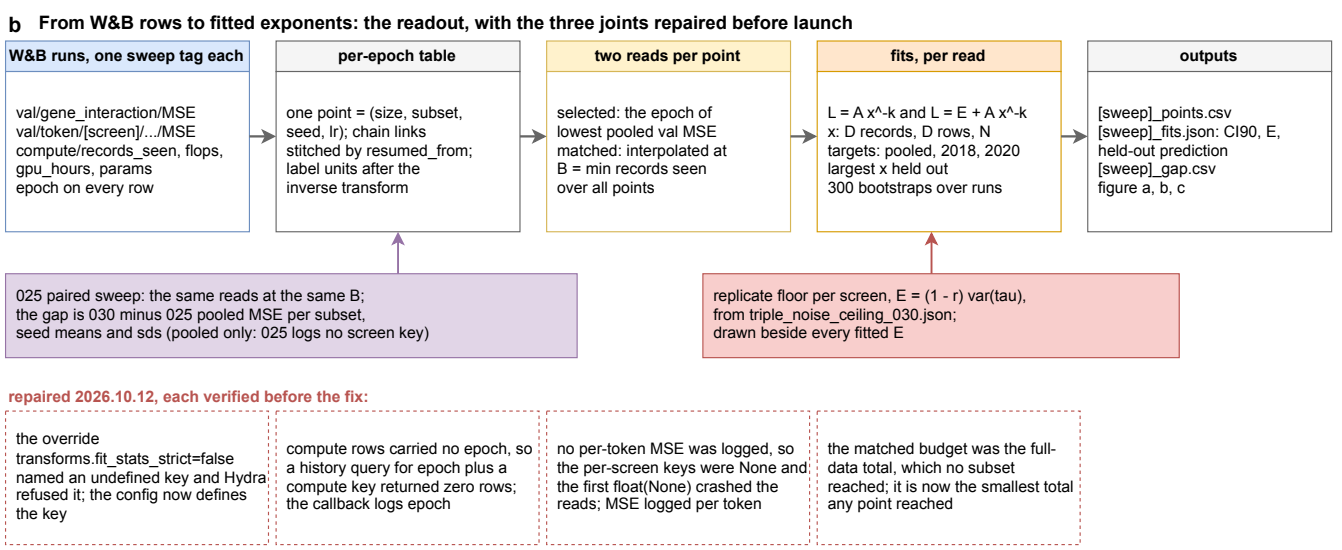
b
From W&B rows to fitted exponents: the readout, with the three joints repaired before launch
W&B runs, one sweep tag each
val/gene_interaction/MSE
val/token/[screen]/.../MSE
compute/records_seen, flops,
gpu_hours, params
epoch on every row
per-epoch table
one point = (size, subset,
seed, lr); chain links
stitched by resumed_from;
label units after the
inverse transform
two reads per point
selected: the epoch of
lowest pooled val MSE
matched: interpolated at
B = min records seen
over all points
fits, per read
L = A x^-k and L = E + A x^-k
x: D records, D rows, N
targets: pooled, 2018, 2020
largest x held out
300 bootstraps over runs
outputs
[sweep]_points.csv
[sweep]_fits.json: CI90, E,
held-out prediction
[sweep]_gap.csv
figure a, b, c
025 paired sweep: the same reads at the same B;
the gap is 030 minus 025 pooled MSE per subset,
seed means and sds (pooled only: 025 logs no screen key)
replicate floor per screen, E = (1 - r) var(tau),
from triple_noise_ceiling_030.json;
drawn beside every fitted E
repaired 2026.10.12, each verified before the fix:
the override
transforms.fit_stats_strict=false
named an undefined key and Hydra
refused it; the config now defines
the key
compute rows carried no epoch, so
a history query for epoch plus a
compute key returned zero rows;
the callback logs epoch
no per-token MSE was logged, so
the per-screen keys were None and
the first float(None) crashed the
reads; MSE logged per token
the matched budget was the full-
data total, which no subset
reached; it is now the smallest total
any point reached