

What eQTL Data Is

And how a meiotic mosaic fits an ontology built for gene edits

Michael Volk

August 27, 2026

Abstract

Expression quantitative trait locus studies cross two strains, collect recombinant progeny, and measure a transcriptome in every one. Four of them rank among the strongest candidates for ingestion, and three are the only datasets in that list that are high-dimensional in both perturbation and readout. They also break an assumption the current schema rests on: the genotype is never observed directly. It is inferred from sparse markers as a mosaic of two parental genomes. This sets out what the records contain, establishes that the strain genotype is nevertheless recoverable to roughly one percent while the causal variant is not recoverable at all, and states what the schema needs before a loader is written.

Section status: ✓ ready ■ author-reviewed, keep stable ✗ not done, agents may edit freely

Provenance: ▲ not in the library mirror; verify at the linked source ● from a review's summary table, not the primary paper

Contents

1	The experiment and its three tables ✗	2
1.1	Two measurements and one inference ✗	2
1.2	Structure worth knowing before modeling it ✗	2
1.3	The four candidates, and why they rank where they do ✗	2
2	The genotype is inferred, not observed ✗	4
2.1	What the methods actually say ✗	4
2.2	Three levels of partial observation ✗	4
2.3	Why sparse observation is nevertheless enough ✗	4
2.4	What reconstruction cannot give you at any marker density ✗	5
2.5	Per cell the answer flips ✗	5
2.6	Reconstruction and attribution come apart ✗	5
3	How it fits the ontology ✗	5
3.1	What already exists ✗	5
3.2	Where the obvious mapping fails ✗	6
3.3	Four gaps, in order of how much they matter ✗	6
3.4	The representation that follows ✗	6
3.5	What to do ✗	7

1 The experiment and its three tables ✗

Two strains that differ genetically are crossed, classically BY, a laboratory strain near-isogenic to S288C, against RM, a vineyard isolate. The F1 hybrid is sporulated and n haploid progeny are collected. Each progeny strain, a **segregant**, is a **meiotic mosaic**: at every locus it carries either the BY or the RM allele, in linked blocks whose boundaries are set by crossover positions. Every segregant’s transcriptome is measured, and the question asked per gene is which loci in the mosaic predict its expression.

The perturbation is therefore not applied to a strain. It is inherited, and it is present everywhere in the genome at once.

1.1 Two measurements and one inference ✗

Genotype matrix. X , segregants by marker positions. Entries encode *parental origin* rather than base identity, so the natural form looks binary. Bloom 2013 used 11,623 markers for 1,008 segregants. BY and RM segregate at roughly 4.5×10^4 sites, so the marker set is a thinned representation of the real variant set, and long runs of constant value along the position axis are what recombination blocks look like. Sec. 2 corrects the binary reading.

Expression matrix. $Y \in \mathbb{R}^{n \times p}$, segregants by genes, with $p \approx 5,000$ to 6,000. This is the phenotype, and it is a vector per strain rather than a scalar.

QTL table. The derived output. For gene g and marker j the standard model is

$$y_{ig} = \mu_g + \beta_{gj} X_{ij} + \varepsilon_{ig}, \quad \varepsilon_{ig} \sim \mathcal{N}(0, \sigma_g^2)$$

fit marker by marker, with significance from a LOD score and a permutation threshold. A row of the table is a gene, a marker interval, an effect size, a LOD and a variance explained.

The first two are measurements. The third is an inference over them, and it changes if the mapping method, the threshold or the marker density changes while the measurement does not. Sec. 3 returns to what that means for storage.

1.2 Structure worth knowing before modeling it ✗

Cis against trans. A **cis** eQTL sits at or near the gene’s own locus and is usually a promoter variant changing that gene’s own transcription. A **trans** eQTL sits elsewhere, typically in a regulator. Cis effects are strong, sparse and roughly one per affected gene; trans effects are weak and numerous.

Hotspots. Trans effects are not spread evenly across loci. A small number of loci each affect hundreds to thousands of transcripts, one regulatory variant propagating outward. The columns of Y are therefore strongly dependent, a few latent factors explain much of the variance, and treating genes as independent is wrong.


1.3 The four candidates, and why they rank where they do ✗

Table 1: The eQTL rows in the candidate list. *Both axes* marks a dataset that is high-dimensional in perturbation and in readout at once, which in that list is true of three datasets in total.


Dataset	Segregants	Readout	Both axes
Albert 2018	1,012	bulk RNA-seq, 1 condition	yes
Boocock 2025	393	one-pot single cell, 2 carbon sources	yes
N’Guessan 2025	~4,500 profiled	single cell	yes
Smith 2008	109	bulk, glucose and ethanol	no

Boocock 2025 draws 27,744 single cells from those 393 segregants, a median of 17 cells per segregant. The extra axis is cells *within* a genotype, which is what allows eQTLs to be mapped for expression noise and for cell-cycle occupancy rather than only for the mean. Smith 2008 is small, and it is the design that demonstrates an eQTL present in one carbon source and absent in the other.

a The cross, and the two things measured on every segregant




b One segregant, one chromosome: from sparse reads to a posterior




Why sparse reads still suffice. The segregant is never sequenced de novo. At each locus it is ASSIGNED to one of two already-sequenced genomes. Meiosis makes ≈90 crossovers, so a genome is ≈10² blocks, and one marker fixes every cataloged variant in its block. The task is locating ≈10² boundaries, not determining ≈10⁴ values.

What no marker density fixes. Boundary intervals leave ≈0.8% of cataloged sites ambiguous. Non-crossover gene conversion (≈46 per meiosis, ≈2 kb) affects more bases than that and is invisible to sparse markers. De novo mutation, aneuploidy and non-reference sequence sit outside the mosaic model entirely.

c What the two matrices look like



d The resulting map: cis on the diagonal, trans in bands



e What that means for a record

Stored. The pair (genotype, environment) → phenotype, one record per segregant per condition. The phenotype side needs nothing new: RNASeqExpressionPhenotype and PseudobulkExpressionPhenotype already exist. The genotype side should hold the MOSAIC, ≈10¹-10² intervals of (chr, start, end, parent, p) against two hash-pinned parent assemblies, rather than ≈10⁴ variants. That stores the posterior as a posterior, is lossless, and keeps linkage explicit.

Not stored, and one field with nowhere to go. The QTL table is an estimate conditional on method, threshold and marker density, with no source value to check, so it is not a Phenotype. And GenePerturbation REQUIRES systematic_gene_name: a cis-eQTL is usually an intergenic promoter variant with no gene to attach to. The same gap blocks every reporter assay, so the two are one decision.

Figure 1: An eQTL study, where inference enters it, and what it costs. **a**, The cross. Both parents are deep-sequenced, so every variant and both its alleles are known before any segregant is examined. The segregants are then assayed twice, sparsely for genotype and densely for expression, and that asymmetry drives everything else. **b**, One segregant along one chromosome, at four levels. True variant sites are dense and none is read directly; marker calls are sparse and are read counts rather than certainties; the HMM output is a posterior whose ramps are the uncertainty, since a crossover is placed only to within one marker interval; the inferred blocks follow. A non-crossover gene conversion tract sitting between markers produces no ramp at all, because no read reports it. **c**, The two matrices. X shows haplotype blocks rather than noise, so its columns are linked; Y is a vector per strain rather than a scalar. **d**, The map that results. Cis effects fall on the diagonal, each variant acting on its own gene and usually through its promoter; trans hotspots appear as vertical bands where one regulatory variant moves hundreds of transcripts. **e**, The consequence for a record, developed in Sec. 3. Panels **b** and **d** are schematic rather than plotted from data. Composed from the sources in Sec. 1 rather than redrawn from a published figure; source notes/assets/drawio/eqtl-experiment-and-genotype-inference.drawio.

2 The genotype is inferred, not observed ✗

2.1 What the methods actually say ✗

Verbatim from the mirrored Boocock 2025 paper:

“We used a hidden Markov model (HMM) to reconstruct the patterns of inheritance of parental alleles in each cell based on the observed genotypes at expressed SNPs.”

“We defined prior genotype probabilities as 0.5 for each parental variant. Emission probabilities were calculated as previously described for low coverage sequencing data.”

“Genotype probabilities were standardized, and markers in very high LD ($r > 0.999$) were pruned. This LD pruning is approximately equivalent to using markers spaced 4 centimorgans (cm) apart.”

So the genotype matrix is not binary. It is

$$X \in [0, 1]^{n \times m}, \quad X_{ij} = \Pr(\text{segregant } i \text{ carries the RM allele at } j \mid \text{reads})$$

and the eQTL model is fit on the standardized posterior directly. Binarizing at 0.5 is done only to compare cell identities.

2.2 Three levels of partial observation ✗

Table 2: What is actually observed, by study. The single-cell design is the most severely partial: genotype coverage is conditioned on the transcript being expressed, so the genotype observation and the phenotype measurement come from the same reads.

Study	Observed	Genotype is
Bloom 2013	low-coverage WGS of each segregant	calls at $\sim 11,623$ markers
Nguyen Ba 2022	low-coverage WGS plus barcode	calls at markers, $n \sim 10^5$
Boocock 2025	only SNPs that happen to be expressed	an HMM posterior per marker

A gene that is switched off in a cell contributes no genotype information about its own locus. That coupling has no analog in any other dataset in the collection.

2.3 Why sparse observation is nevertheless enough ✗

The segregant is not being sequenced de novo. It is being *assigned* to one of two known genomes at each locus, which is a far easier problem:

1. The parents are deep-sequenced, “deep ($> 100\times$) paired-end sequenced data for our parental strains,” so the complete variant catalog and both alleles at every site are known before any segregant is looked at.
2. Meiosis produces few boundaries. In *S. cerevisiae* roughly 90 crossovers per meiosis, so across 16 chromosomes a segregant genome is on the order of 10^2 contiguous blocks.
3. Within a block, parental origin is constant, so one marker fixes every cataloged variant in its block.

Reconstruction reduces to locating $\sim 10^2$ boundaries rather than determining $\sim 10^4$ independent values.

Error budget, derived from published parameters rather than measured. BY and RM segregate at $\sim 4.5 \times 10^4$ SNPs over 12 Mb, about 1 per 260 bp. Bloom 2013’s 11,623 markers give ~ 1 kb spacing, so a crossover is localized to about 4 variants:

$$\underbrace{90}_{\text{crossovers}} \times \underbrace{4}_{\text{variants per interval}} \approx 360 \text{ of } 4.5 \times 10^4 \approx 0.8\% \text{ ambiguous}$$

leaving roughly 99% of cataloged variant sites assignable. ▲ [external: arithmetic from published crossover and variant-density figures, not a measured concordance]

The larger error source is not crossovers. Yeast meiosis also produces ~ 46 non-crossover gene conversion events per meiosis with tracts of order 2 kb. Those are short, they do not shift the surrounding haplotype, and sparse markers miss them entirely, so they affect more base pairs than boundary uncertainty and are harder to detect.

2.4 What reconstruction cannot give you at any marker density ✗

- **De novo mutations.** A segregant is a mosaic plus whatever arose since the cross. Few per strain, invisible to a parental-assignment method, so the genome is never exactly a mosaic.
- **Aneuploidy and copy number.** These need coverage, not markers, and yeast segregants do acquire them.
- **Non-reference sequence.** A variant call file is not an assembly. Reconstructing total genomic content needs the parents' assemblies, including sequence absent from S288C.

2.5 Per cell the answer flips ✗

Boockock reports *median genotype agreement of 92.5%* between the single-cell HMM genotypes and whole-genome sequencing of the same strains, improving with UMI count. About 1 marker in 13 wrong is entirely acceptable for an association fit across 27,744 cells, where the posterior enters as a standardized covariate and the error averages out. It is nowhere near sufficient to write down a genome. What rescues the strain-level genotype is pooling: a median of 17 cells per segregant, against 393 segregants that had already been whole-genome sequenced in Bloom 2013 and against which the cells were matched.

Do not conflate reconstruction resolution with mapping resolution. The HMM used all expressed SNPs. The 4 cM LD-pruned marker set was used for the association, because markers above $r > 0.999$ carry no independent information. Coarse mapping resolution says nothing about how well the genome can be reconstructed.

2.6 Reconstruction and attribution come apart ✗

which allele is at a site ✓ against which variant causes the effect ✗

The first is a two-way assignment with strong priors and succeeds at $\sim 99\%$. The second is blocked by collinearity: within a haplotype block every variant is perfectly correlated with every other, so an effect is identifiable only up to the block, and no number of additional segregants changes that. Pruning markers above $r > 0.999$ concedes exactly this point.

The consequence for the candidate list is favorable. The sequence gate asks whether the total genomic content of a strain can be estimated, not whether a causal variant can be identified, so **segregant panels pass**, conditional on holding the parent assemblies. The consequence for modeling is that a cross localizes an effect to an interval and nothing more. Only an assay that tests alleles one at a time, such as a massively parallel reporter assay, gets from an interval to a base, which is why those datasets belong in the same ingestion wave rather than a later one.

3 How it fits the ontology ✗

3.1 What already exists ✗

The primary record is a genotype and an environment mapped to a phenotype, which is exactly **Experiment**. Nothing conceptual is missing, and four pieces are already built.

- **RNASeqExpressionPhenotype** is the bulk readout, and it was written for this situation. Its docstring describes a survey “where each isolate’s transcriptome is measured on its OWN genome,” with “no common reference to ratio against,” storing absolute TPM. Albert 2018 and Smith 2008 fit unchanged.
- **PseudobulkExpressionPhenotype** is the single-cell readout, built for Nadal-Ribelles 2025. It carries **dispersion** and **n_cells** per genotype, which is the right shape for Boockock’s noise eQTLs: heterogeneity survives as a per-genotype scalar rather than being averaged away.
- **SequenceVariantPerturbation** states the intent outright. Together with **CopyNumberVariantPerturbation** it “lets a NATURAL ISOLATE be modeled as a perturbation set off S288C ($\sim 4,500$ sequence variants per isolate).”
- The off-graph pointer pattern on that class, **sequence_source**, **sequence_uri** and **sequence_sha256**, means variant sequences are never inlined.

3.2 Where the obvious mapping fails ✗

Reusing `SequenceVariantPerturbation` for a segregant is the obvious move and it is wrong. **A natural isolate and a segregant are not the same kind of object.** Caudal's isolates are each individually sequenced, so the variant set is observed and every allele dereferences to a real sequence with a hash. A segregant is never sequenced to that standard; the *parents* are, and the segregant is imputed as a mosaic of them. That class promises a dereferenceable, hash-verified sequence, and for a segregant no such artifact exists. Using it would encode an inference as an observation, which is the precise failure the verification levels exist to prevent.

3.3 Four gaps, in order of how much they matter ✗

Intergenic variants have no representation. `GenePerturbation` requires `systematic\_gene\_name` and validates it against

```
^(Y[A-P][LR]\\d{3}[WC](-[A-Z])?)|Q\\d{4}|YNC[A-Q]\\d{4}[WC])\\$
```

so every perturbation must attach to a real gene. A promoter SNP 300 bp upstream of an ORF has no gene to attach to, and assigning it to the downstream gene is an interpretation, wrong whenever the variant acts on a divergent neighbor. This is the sharpest gap and it is not incidental: the promoter is where the strongest and most interpretable cis signal lives. It is also the same gap that blocks every massively parallel reporter assay, which perturbs regulatory sequence exclusively.

Cardinality. An SGA record carries at most three perturbations. A segregant carries $\sim 10^4$. For $n \sim 10^3$ segregants that is $\sim 10^7$ perturbation objects for one dataset, and `Genotype` sorts and set-compares its perturbation list on every equality check while the L1 key hashes the genotype signature.

Reference asymmetry. A segregant's natural reference is one of its two parents, but the convention expresses everything off S288C. BY is near-isogenic to S288C, so BY-derived blocks are nearly empty perturbation sets while RM-derived blocks carry the full variant load. Keeping S288C is still right, since it is what makes these records comparable to everything else, but the lopsidedness reflects the reference choice rather than the biology and should be documented rather than discovered.

Linkage is destroyed by a set. Variants arrive in blocks and co-segregate. Representing a genotype as an unordered set discards exactly the structure that makes mapping possible and that limits attribution.

3.4 The representation that follows ✗

Store what was actually determined, a **haplotype mosaic** rather than a variant set:

$$G_i = \{(\text{chr}, \text{start}, \text{end}, \text{parent}, p)\}, \quad |G_i| \sim 10^1 - 10^2$$

against two hash-pinned parent assemblies, with a per-interval posterior p and the marker spacing recorded as the resolution. This is better on four counts at once.

1. **Honest.** It stores the inference as an inference, with its uncertainty, rather than promising sequence that was never read.
2. **Compact.** Tens of intervals per strain instead of $\sim 10^4$ variants, which dissolves the cardinality problem rather than solving it.
3. **Lossless.** The full variant set is recoverable exactly from the two parent assemblies plus the intervals, so it becomes a derived view and nothing is given up.
4. **Linkage-preserving.** The interval is the unit, so co-segregation is explicit.

The cost is a genuinely new perturbation type, keyed on a genomic interval and a parent rather than on a gene, carrying a probability. That is the same non-gene-keyed requirement the promoter-variant gap already imposes, so the two should be resolved together.

3.5 What to do x

1. **Ingest the two matrices, not the QTL table.** One `RNASeqExpressionExperiment` per segregant per condition, or `PseudobulkExpressionExperiment` for the single-cell studies. The QTL table is an estimate conditional on method and threshold, it has no source value to check against at any verification level, and storing it as a **Phenotype** would be ingesting somebody's analysis as data.
2. **Decide the intergenic question deliberately**, because it gates both this class and the regulatory-DNA class. Either a non-gene-keyed perturbation leaf on a genomic interval, or gene-keying with the attribution recorded as a **ProvenanceGap**, or restriction to coding variants. The third is the honest short-term option and the worst long-term one.
3. **Treat the parent assemblies as an ingestion dependency.** A segregant record means nothing without them, so they are part of the artifact rather than context, and each needs a hash-pinned assembly rather than a variant call file.
4. **Measure the cardinality cost** on a single Albert 2018 segregant before building anything.