

# Accelerating metabolic engineering through multimodal phenotype prediction with TorchCell and the Cell Graph Transformer

Michael Volk<sup>1\*</sup>, First Coauthor<sup>1</sup> and Huimin Zhao<sup>1\*</sup>

<sup>1</sup>Department of Chemical and Biomolecular Engineering, University of Illinois Urbana-Champaign, Urbana, IL, USA.

\*Corresponding author(s). E-mail(s): [mjvolk3@illinois.edu](mailto:mjvolk3@illinois.edu);  
[zhao5@illinois.edu](mailto:zhao5@illinois.edu);

Contributing authors: ;

## Abstract

Deep learning has advanced protein structure prediction, yet metabolic engineering lags because phenotype data are scattered across heterogeneous studies lacking a shared ontology. We present TorchCell, an open-source framework that annotates literature and public datasets with a shared ontology and ingests them into a Neo4j knowledge graph, then builds deep-learning-ready datasets in which each strain is a perturbation operator over a shared wildtype reference cell combining genome sequence, gene networks, metabolism, and environment. We develop the Cell Graph Transformer (CGT), a virtual-cell architecture in which biological interaction graphs constrain multi-head attention and an equivariant perturbation operator feeds multitask decoders. In *Saccharomyces cerevisiae*, CGT predicts trigenic gene-gene interactions (Pearson  $r = \mathbf{0.454}$ ), outperforming DANGO, DCell, and Yeast9 flux balance analysis, and generalizes to knockout expression ( $r = \mathbf{0.543}$ ) and morphology ( $r = \mathbf{0.619}$ ). CGT recommends gene deletions for beta-carotene and betaxanthin production, pairing with autonomous biofoundries for AI-guided strain engineering.

**Keywords:** metabolic engineering, foundation models, gene interactions, virtual cell, knowledge graph, strain design

# 1 Introduction

*[FILLER – replace with Intro prose.]* Deep learning transformed protein structure and function prediction, yet systems biology and metabolic engineering lag because phenotype data are scattered across small, heterogeneous studies with no shared schema [1]. Prior approaches fall short: mechanistic flux balance analysis misses epistasis, and recent perturbation transformers do not consistently beat linear baselines [2]. Our own classical baselines predict knockout fitness but fail on gene interactions (Suppl. Fig. S1), motivating a deep, biologically grounded model. Here we present TorchCell and the Cell Graph Transformer (CGT), whose data model and architecture we outline in Fig. 1.

## 2 Results

### **A shared ontology and knowledge graph unify metabolic-engineering data.**

*[FILLER – R1.]* TorchCell annotates literature and public datasets with a shared experiment ontology and ingests them into a Neo4j knowledge graph; dataloaders emit deep-learning-ready datasets in which each strain is a perturbation operator over a shared wildtype reference cell (Fig. 1; data scale in Suppl. Fig. S2; ontology/KG spec in Suppl. Fig. S4).

### **The Cell Graph Transformer: a graph-constrained virtual cell.**

*[FILLER – R2.]* Biological interaction graphs constrain multi-head attention through a loss aligning attention to network priors; an equivariant perturbation operator maps the reference embedding to the perturbed strain state (Fig. 2; ablations in Suppl. Fig. S3).

### **CGT predicts trigenic gene-gene interactions at state of the art.**

*[FILLER – R3.]* Classical ML fails on this task (Suppl. Fig. S1); CGT reaches Pearson  $r = 0.454$ , outperforming DANGO, DCell, and Yeast9 FBA (Fig. 3).

### **One embedding, many phenotypes: expression and morphology.**

*[FILLER – R4.]* The same architecture predicts knockout expression ( $r = 0.543$ ) and single-knockout morphology ( $r = 0.619$ ) (Fig. 4; expression diagnostics in Suppl. Fig. S5).

### **CGT-guided strain design with an autonomous foundry.**

*[FILLER – R5.]* We recommend gene deletions for beta-carotene and betaxanthin production and pair with the UIUC iBioFoundry for iterative design (Fig. 5; inference at scale in Suppl. Fig. S6).

## 3 Discussion

*[FILLER – Discussion.]* A strain-aware, biologically grounded representation is the differentiator; limitations include single-organism scope, morphology-data maturity, and error-bar provenance.

## 4 Methods

[*FILLER – Methods.*] Ontology and knowledge-graph construction; dataset and perturbation-operator formalism; CGT architecture and the graph-regularized attention loss; training, splits, metrics; baselines (DANGO, DCell, Yeast9 FBA); inference at scale; statistics and error-bar definitions.


**Supplementary information.**

**Acknowledgments.**


## Declarations

## References


- [1] Volk, M. J., Tran, V. G., Tan, S.-I., others & Zhao, H. Metabolic engineering: Methodologies and applications. *Chemical Reviews* **123**, 5521–5570 (2023).
- [2] Ahlmann-Eltze, C., Huber, W. & Anders, S. Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines. *Nature Methods* **22**, 1657–1661 (2025).
- [3] Ma, J. *et al.* Using deep learning to model the hierarchical structure and function of a cell. *Nature Methods* **15**, 290–298 (2018).
- [4] Zhang, R., Ma, J. & Ma, J. Dango: Predicting higher-order genetic interactions. *bioRxiv* (2020).
- [5] Merzbacher, C., Mac Aodha, O. & Oyarzún, D. A. Accurate prediction of gene deletion phenotypes with flux cone learning. *Nature Communications* **16**, 8492 (2025).
- [6] Bunne, C. *et al.* How to build the virtual cell with artificial intelligence: Priorities and opportunities. *Cell* **187**, 7045–7063 (2024).




**Fig. 1** TorchCell: ontology-unified knowledge graph and the perturbation-operator cell representation. Panels (a)–(d).




**Fig. 2** The Cell Graph Transformer architecture. (a)–(d).



**Fig. 3** State-of-the-art trigenic gene-gene interaction prediction. (a)–(d).



**Fig. 4** One latent cell embedding generalizes across phenotype classes. (a)–(d).



**Fig. 5** CGT-guided strain design closes the DBTL loop with an autonomous foundry. (a)–(d).